

Patterns based on Semaphores

CS511

Review of Semaphores

- ▶ An Abstract Data Type with two operations
 - ▶ acquire
 - ▶ release
- ▶ Can be used to solve the mutual exclusion problem
- ▶ Can be used to synchronize cooperative threads

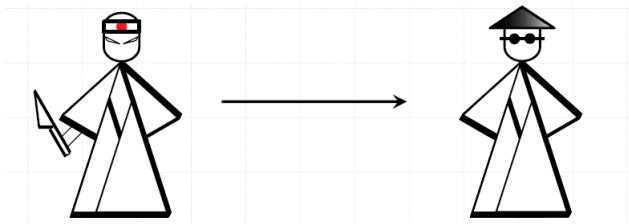
Today

- ▶ Recurring problems in the area
- ▶ Proven solution templates

Producers/Consumers

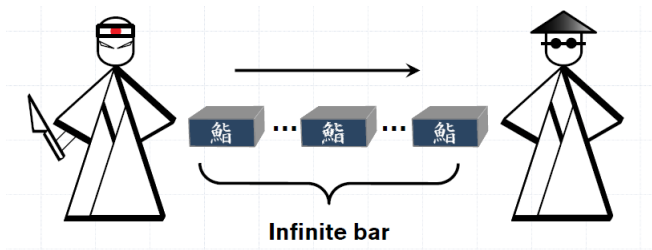
Readers/Writers

Producers/consumers



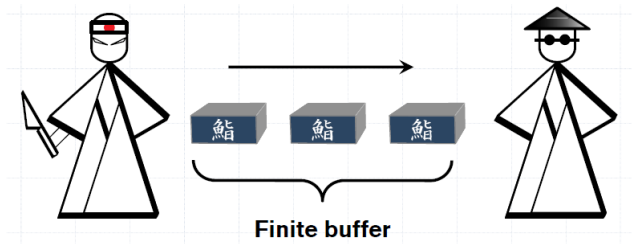
- ▶ A common pattern of interaction
- ▶ Must cater for difference in speed between each party

Unbounded Buffer



- ▶ The producer can work freely
- ▶ The consumer must wait for the producer to produce

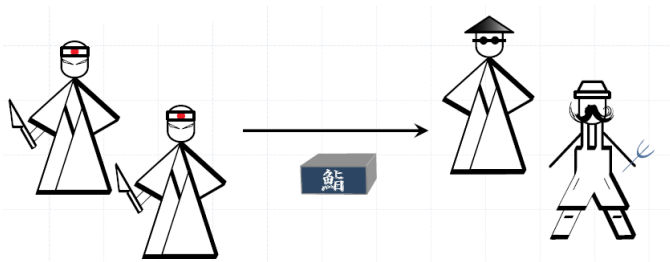
Bounded Buffer



- ▶ The producer must wait when the buffer is full
- ▶ The consumer must wait for the producer to produce

Buffer using Semaphores

- ▶ Capacity 1



- ▶ Various producers
- ▶ Various consumers
- ▶ Semaphores

Buffer using Semaphores – 1 producer and 1 consumer

```
1 Object buffer;
2 ...
3 ...

4 Thread.start { // Prod 4 Thread.start { // Cons
5     while (true) {      5     while (true) {
6         ...              6         ...
7         buffer = produce(); 7         consume(buffer);
8         ...              8         ...
9     }                   9     }
```

Split Binary Semaphores

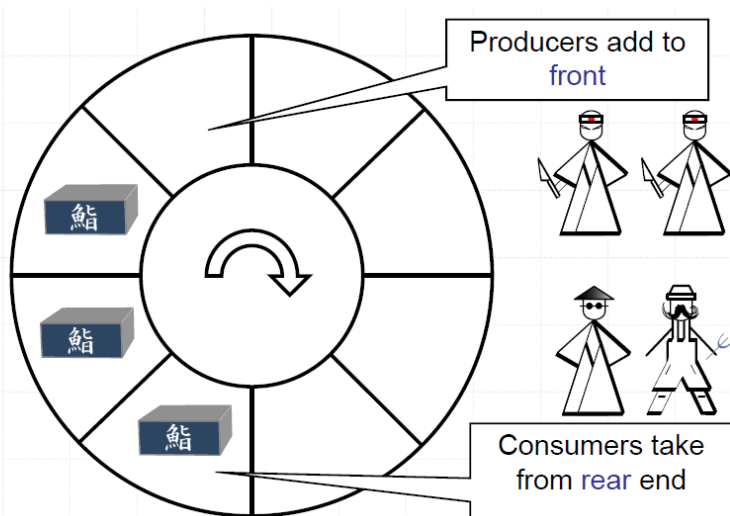
- ▶ Two semaphores
 - ▶ one to hold permissions to produce
 - ▶ one to hold permissions to consume
- ▶ Initialization
 - ▶ `produce = 1`
 - ▶ `consume = 0`
- ▶ Invariant
 - ▶ `produce + consume <= 1`

Split Binary Semaphores

```
1 Object buffer;
2 Semaphore produce = new Semaphore(1);
3 Semaphore consume = new Semaphore(0);

4 Thread.start { // Prod 4 Thread.start { // Cons
5     while (true) {      5     while (true) {
6         produce.acquire(); 6         consume.acquire();
7         buffer = produce(); 7         consume(buffer);
8         consume.release(); 8         produce.release();
9     }                   9     }
```

N Size Buffer



General Semaphores

- ▶ Semaphores count the number of empty slots in the buffer
- ▶ Initialization
 - ▶ There are N empty slots
 - ▶ There are 0 full slots
- ▶ Invariant
 - ▶ $\text{produce} + \text{consume} \leq N$

Unique Producer/Consumer

```
1 Object[]  buffer  = new Object[N];
2 Semaphore produce = new Semaphore(N);
3 Semaphore consume = new Semaphore(0);
4 int start = 0;
5 int end    = 0;

6 Thread.start { // Prod          6 Thread.start{ // Cons
7     while (true) {              7     while (true) {
8         ...                      8         ...
9         buffer[start] = produce(); 9         consume(buffer[end]);
10        start = (start+1) % N;      10        end = (end+1) % N;
11        ...                       11        ...
12    }                             12    }
```

Unique Producer/Consumer

```
1 Object[]  buffer = new Object[N];
2 Semaphore produce = new Semaphore(N);
3 Semaphore consume = new Semaphore(0);
4 int start = 0;
5 int end = 0;

6 Thread.start{ // Prod          6 Thread.start { // Cons
7     while (true) {             7     while (true) {
8         produce.acquire();      8         consume.acquire();
9         buffer[start] = produce(); 9         consume(buffer[end]);
10        start = (start+1) % N;    10        end = (end+1) % N;
11        consume.release();       11        produce.release();
12    }                           12    }
```

Multiple Producers

- ▶ We cannot simply add multiple instances of the producer
- ▶ Why? Justify with a trace
- ▶ What can we do about it?

```
1 // declarations: same as above...
```

6	Thread.start{ // ProdA	6	Thread.start { // ProdB
7	while (true) {	7	while (true) {
8	produce.acquire();	8	produce.acquire();
9	buffer[start] = produce();	9	buffer[start] = produce();
10	start = (start+1) % N;	10	start = (start+1) % N;
11	consume.release();	11	consume.release();
12	}	12	}

Multiple Producers

- ▶ Must guarantee mutual exclusion between producers:
 - ▶ We add a new semaphore

```
1 Semaphore mutexP = new Semaphore(1);
2
3 Thread.start { // Prod
4     while (true) {
5         produce.acquire();
6         mutexP.acquire();
7         buffer[start] = produce();
8         start = (start+1) % N;
9         mutexP.release();
10        consume.release();
11    }
12 }
```

Multiple Consumers

- ▶ Must guarantee mutual exclusion between consumers

```
1 Semaphore mutexC = new Semaphore(1);
2
3 Thread.start { // Cons
4     while (true) {
5         consume.acquire();
6         mutexC.acquire();
7         consume(buffer[end]);
8         end = (end+1) % N;
9         mutexC.release();
10        produce.release();
11    }
12 }
```

Putting it all together

```
1  int N = 10;
2  int[] buffer = new int[N];
3  Semaphore produce = new Semaphore(N);
4  Semaphore consume = new Semaphore(0);
5  int start = 0;
6  int end = 0;
7  int counter = 0;
8
9  void consume(int i) { }
10
11 int produce () { return counter++; }
12
13 Semaphore mutexP = new Semaphore(1);
14 Semaphore mutexC = new Semaphore(1);
15
16 // continues in next slide
```

Putting it all together

```
1 Consumer(int id) {
2     while (true) {
3         consume.acquire();
4         mutexC.acquire();
5         consume(buffer[end]);
6         print(id+" consumed product "+ buffer[end] + " at "+ end);
7         end = (end+1) % N;
8         mutexC.release();
9         produce.release();
10    }
11 }
12
13
14 Producer(int id) {
15     while (true) {
16         produce.acquire();
17         mutexP.acquire();
18         buffer[start] = produce();
19         print(id+" add product "+ buffer[start]+ " at "+ start);
20         start = (start+1) % N;
21         mutexP.release();
22         consume.release();
23    }
24 }
```

Putting it all together

```
1
2 for (i=0; i<5; i++) {
3     int id = i;
4     thread Producer(id);
5     thread Consumer(id);
6 }
```

Producers/Consumers

Readers/Writers

Readers/Writers

- ▶ There are shared resources between two types of threads
 - ▶ **Readers:** access the resource without modifying it
 - ▶ **Writers:** access the resource and may modify it
- ▶ Mutual exclusion is too restrictive
 - ▶ **Readers:** can access simultaneously
 - ▶ **Writers:** at most one at any given time

Properties a Solution should Possess

- ▶ Each read/write operation should occur inside the critical region
- ▶ Must guarantee mutual exclusion between the writers
- ▶ Must allow multiple readers to execute inside the critical region simultaneously

First Solution: Priority Readers

```
1 Writer() {  
2  
3     ...  
4     write();  
5     ...  
6  
7 }
```

```
1 Reader() {  
2  
3     ...  
4     read();  
5     ...  
6  
7 }
```

First Solution: Priority to Readers

- ▶ One semaphore for controlling write access
- ▶ Before writing, the permission must be obtained and then released when done
- ▶ The first reader must “steal” the permission to write and the last one must return it
 - ▶ We must count the number of readers inside the CS
 - ▶ This must be done inside its own CS

First Solution: Priority Readers

```
1 Semaphore resource = new Semaphore(1);
2 Semaphore numReadersMutex = new Semaphore(1);
3 int numReaders = 0;

1 Writer() {
2     resource.acquire();
3     write();
4     resource.release();
5 }

1 Reader() {
2     numReadersMutex.acquire();
3     numReaders++;
4     if (numReaders == 1)
5         resource.acquire();
6     numReadersMutex.release();
7
8     read();
9
10    numReadersMutex.acquire();
11    numReaders--;
12    if (numReaders == 0)
13        resource.release();
14    numReadersMutex.release();
15 }
```

Note: Is this solution free from starvation?

Second Solution: Priority Writers

- ▶ The readers can potentially lock out all the writers
 - ▶ We need to count the number of writers that are waiting
 - ▶ Also, this counter requires its own CS
- ▶ Before reading the readers must obtain a permission to do so

Second Solution: Priority Writers

```
1 Writer() {
2     numWritersMutex.acquire();
3     numWriters++;
4     if (numWriters == 1)
5         readTry.acquire();
6     numWritersMutex.release();
7
8     resource.acquire();
9     write();
10    resource.release();
11
12    numWritersMutex.acquire();
13    numWriters--;
14    if (numWriters == 0)
15        readTry.release();
16    numWritersMutex.release();
17 }

1 Reader() {
2     readTry.acquire();
3     numReadersMutex.acquire();
4     numReaders++;
5     if (numReaders == 1)
6         resource.acquire();
7     numReadersMutex.release();
8     readTry.release();
9
10    read();
11
12    numReadersMutex.acquire();
13    numReaders--;
14    if (numReaders == 0)
15        resource.release();
16    numReadersMutex.release();
17 }
```

- ▶ Readers might starve
- ▶ Solution in which neither readers nor writers starve?
 - ▶ Hint: Common service queue for both readers and writers

Third Solution

```
1 int numReaders;
2 Semaphore resource = new Semaphore(1);
3 Semaphore readCountAccess = new Semaphore(1);
4 Semaphore serviceQueue = new Semaphore(1);

1 Writer() {
2
3
4     serviceQueue.acquire();
5     resource.acquire();
6     serviceQueue.release();
7
8
9     writeResource();
10
11    resource.release();
12 }

1 Reader() {
2     serviceQueue.acquire();
3     readCountAccess.acquire();
4     readCount++;
5     if (readCount == 1)
6         resource.acquire();
7     readCountAccess.release();
8     serviceQueue.release();
9
10    readResource();
11
12    readCountAccess.acquire();
13    readCount--;
14    if (readCount == 0)
15        resource.release();
16    readCountAccess.release();
17 }
```

Summary

1. Semaphores are elegant and efficient for solving problems in concurrent programs
2. Still, they are low-level constructs since they are not structured
3. Monitors will provide synchronization by encapsulation